



TriageLine: an assistant you can talk over

- Theme 05 asks for a real-time agent that keeps working while the caller interrupts, corrects and changes their mind.
- TriageLine answers immediately, runs tools in the background, and throws away anything the caller has made obsolete.



Architecture

- LiveKit room;
- Silero VAD; hosted STT (Gemini, Deepgram or Groq).
- ParticipantAgent: a fast path that acknowledges within milliseconds and a slow path for tools, ASR and vision.
- Epoch-guarded coordination and an idempotency ledger.
- Hosted TTS back into the room.



FDB-v3 benchmark

- Official Full-Duplex-Bench v3 at pinned commit 3e799c45 via `./run_fdb_v3.sh`, which fails loudly on runner errors or stale results.
- Live judged pass rate: [pending — run `./run_fdb_v3.sh --require-judge`].
- First-response latency: [from results/results.md].
- The offline text replay (85/100, judge off) is a diagnostic only.



Interruption and recovery

- Every correction bumps an epoch.
- Read-only work is cancelled and stale results are dropped.
- State changes are never guessed: they reconcile through the ledger.
- "Actually, don't do that" in the same breath performs nothing.



Honest evaluation

- `scripts/integrity_audit.py` checks decision code for benchmark phrasing (PASS).
- An independent held-out set of 30 new phrasings scores 26/30 with rules only.
- The LLM planner (Gemini, with Cerebras, OpenRouter and Mistral failover) is schema-validated and cannot skip confirmation.



Extension: Triage Line

- Roadside-incident call flow on the same LiveKit shell: dialogue, deliberation, then a commit state machine (proposed, pending confirmation, finalized or aborted).
- A mid-prompt correction aborts the stale action; nothing is dispatched without an explicit yes, and every dispatch is labelled simulated.



Product surface

- Mobile PWA with text, voice clips and camera frames.
- Continuous WebRTC voice at /rtc.html.
- Authenticated token API for native LiveKit SDKs.
- One ASGI server, Docker and compose.
- All business tools are simulated.



Results and next steps

- 210+ Python and 10 JS tests pass; the practice harness scores 89/100.
- Still to prove: the judged live FDB-v3 run with team keys, real-room barge-in checks, real provider adapters, user identity and persistent sessions.